Tan Dai Ngo

250-661-9530 | Victoria, BC | ngotandai95@gmail.com | linkedin.com/in/ntdai95/ | ntdai95.github.io

Technical Skills

- **Programming:** Python, Java, SQL, MATLAB
- **Machine Learning:** scikit-learn, XGBoost, CatBoost, regression, classification, clustering, feature engineering, Optuna
- **Data & ML Systems:** Apache Spark, Kafka, Cassandra, Neo4j, MLflow, ETL pipelines, data processing, distributed systems
- **Deployment & GenAI:** FastAPI, Docker, AWS (EC2), Sentence Transformers, Qdrant, Retrieval-Augmented Generation (RAG), Ollama LLM

Education

Master of Engineering in Applied Data Science — **September 2025 - December 2027 (Expected)**

University of Victoria

- **GPA:** 9.0 / 9.0
- **Relevant Coursework:** Optimization for Machine Learning, Data Analysis and Pattern Recognition, Systems for Massive Datasets, Information Theory, Algorithms and Data Models, Data Mining, Applied Data Analysis

Master of Science in Computer Science — University of Chicago

Bachelor of Science in Economics — University of Washington

Experience

Software Developer | T-Mobile (via BeaconFire Inc.) — *Bellevue, WA, USA* | **June 2022 - April 2024**

- Improved reliability of operational data workflows by automating Kafka data reprocessing pipelines, reducing manual intervention by 80%.
- Reduced data retrieval latency for partner workflows to under 3 seconds by redesigning Cassandra schemas supporting 100k+ records.
- Enabled scalable processing of roaming and partner datasets by developing microservices supporting downstream analytics and operational systems.

Projects

End-to-End Machine Learning System for Multi-Output Fuel Blending — Shell.ai Hackathon 2025

- Deployed a production ready machine learning system using FastAPI, Docker, and AWS EC2 to serve real time multi-output regression predictions for 10 chemical blend properties.
- Benchmarked XGBoost vs. CatBoost using 5-fold cross-validation, selecting CatBoost after achieving lower error (MAPE 0.64 vs. 1.29) and productionizing the best performing model.
- Engineered robust feature and inference pipelines using weighted property aggregation, entropy based mixture metrics, and serialized preprocessing to ensure consistent production inference.

Distributed Ocean Data ML Platform with RAG — Systems for Massive Datasets Project

- Built a distributed ML pipeline using Apache Spark (Bronze → Silver → Gold architecture) to ingest and harmonize 10M+ NOAA and ONC ocean sensor observations for scalable feature engineering and time series forecasting.
- Trained and evaluated forecasting models using XGBoost with Optuna hyperparameter tuning and MLflow experiment tracking, applying holdout validation based on time to ensure reliable performance on unseen temporal data.
- Developed a FastAPI RAG service using Ollama LLM, Sentence Transformers, and Qdrant vector search, implementing retrieval benchmarking (hit@k = 1.0, term recall = 0.875).

Stock Sentiment Analysis (Team of 4) — Algorithms and Data Models Project

- Built a sentiment analysis pipeline over 15k+ tweets across 10 tech stocks, enabling large-scale market sentiment analysis, by integrating X API data, SQLite/Neo4j storage, and FinTwitBERT modeling.
- Identified measurable sentiment-return relationships, including statistically significant predictive effects for NVIDIA, by implementing engagement-weighted sentiment aggregation and Granger causality testing.
- Analyzed equity volatility dynamics across major tech stocks, revealing persistent volatility clustering, by applying GARCH(1,1) modeling and automated cross-stock analysis workflows.